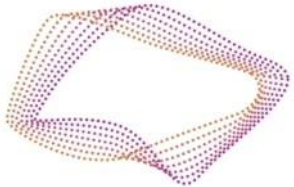

Classifying Multimodal Data Using Transformers

Watson Chua, Lu Li, Alvina Goh, Amelia Lee
Data Science and Artificial Intelligence Division
Government Technology Agency of Singapore

Background

- GovTech successfully built a transformers-based classifier which can automatically predict the government agencies to handle feedback on municipal issues from residents, using the following:
  - Text
  - Geolocation
  - Images

- With this experience, the team is conducting a hands-on tutorial to flatten the learning curve for data scientists and machine learning engineers who want to apply machine learning on multimodal data.

Objective

- At the end of the tutorial, you will be able to use PyTorch and Hugging Face Transformers to build:
  - A text classifier using BERT
  - A text-image classifier with dual encoders using BERT and ResNet
  - A text-image classifier with joint encoders using ALBEF

- We will be using a selected subset of the Webvision1.0 dataset for the tutorial.

Tutorial Outline

No.	Item	Duration
1	Sharing on the Municipal Issues Feedback Classifier	35 mins
2	Building a Text Classifier with BERT	35 mins
3	Building a Dual-Encoder Text-Image Classifier with ResNet and BERT	35 mins
4	Building a Joint-Encoder Text-Image Classifier with ALBEF	35 mins
5	Discussion/ Q&A	15 mins

20 mins break at 3.30pm

Sharing on the Municipal Issues Feedback Classifier

Municipal Services Office (MSO)

- In charge of feedback management and service delivery for municipal services
- Receives feedback on municipal issues from the public and channels them to other government agencies and town councils to resolve the issues

Municipal Issues Feedback

Pigeons at hawker centres

Overgrown greenery

Fallen lamppost

Feedback is received from 3 channels:

Manual Case Routing Process

Text
Image
Geolocation

MSO receives a feedback case

An officer selects an agency to the route case to

Agency handles feedback

MSO is updated when the case is closed

Case ID	Description	Images	X-Coord	Y-Coord	Agency
1	Potholes on the road	…	…	…	Land Transport Authority
2	Fallen tree in the park	…	…	…	National Parks
3	Mosquitoes breeding in my estate	…	…	…	National Environmental Agency
…	…	…	…	…	…

417k cases handled each year!

Predicting Agency to Handle Feedback

New Feedback

The recently patched road has a new pothole appearing. Very obvious.

1.3404329
103.6901702

Text

Image

Geolocation

App or Chatbot

Classifier trained on historical data from CRM

Agency Classifier

Prediction

Land Transport Authority

Training the Classifier: Text Encoder + Geolocs Architecture

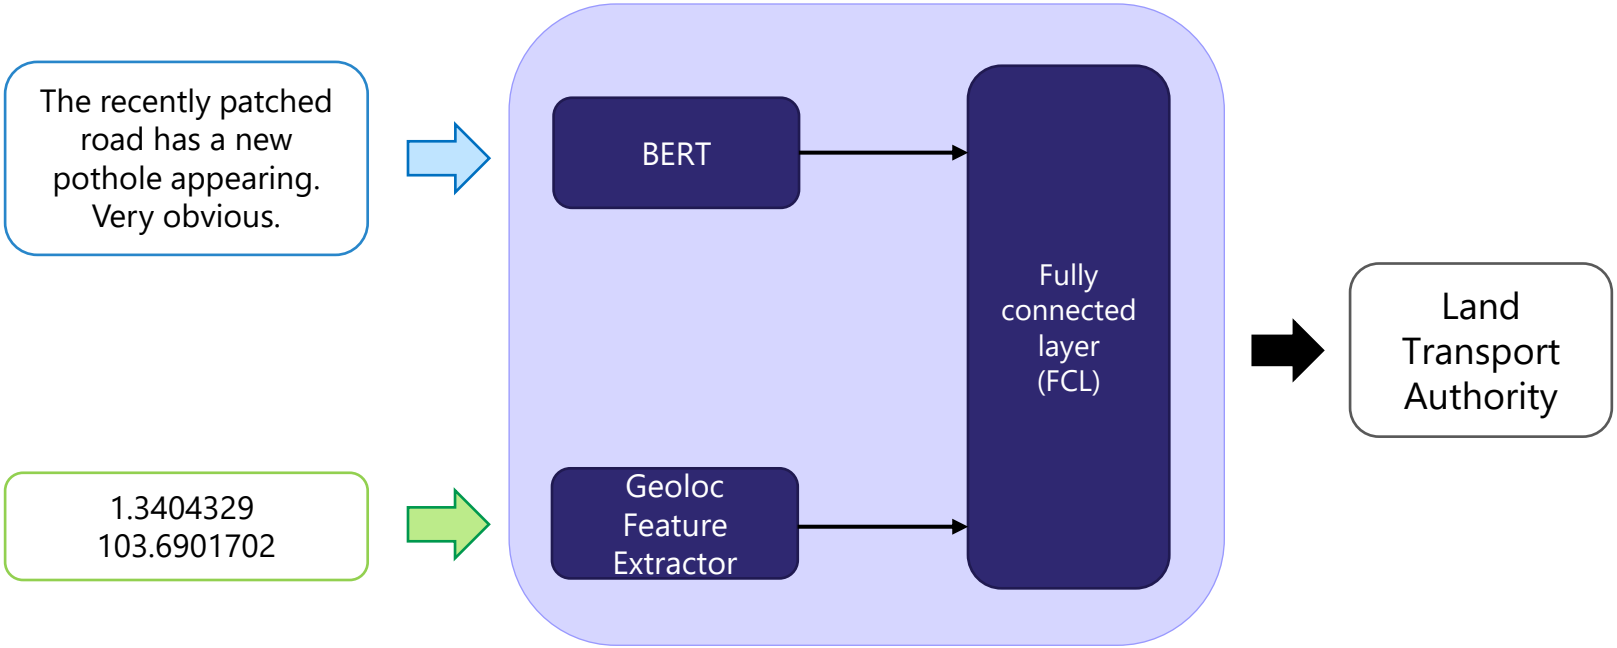

Training the Classifier: Dual-Encoder Text-Image + Geolocs Architecture

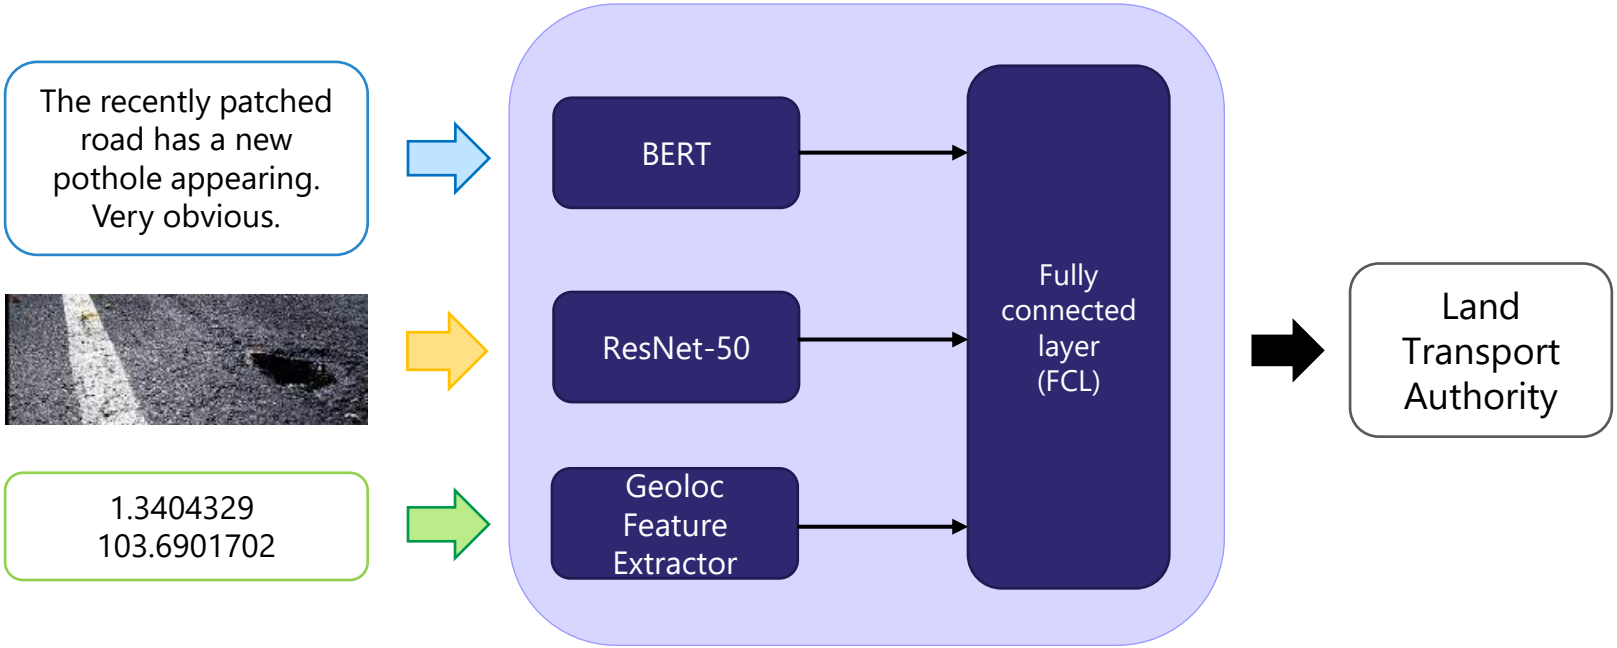

Training the Classifier: Joint-Encoder Text-Image + Geolocs Architecture

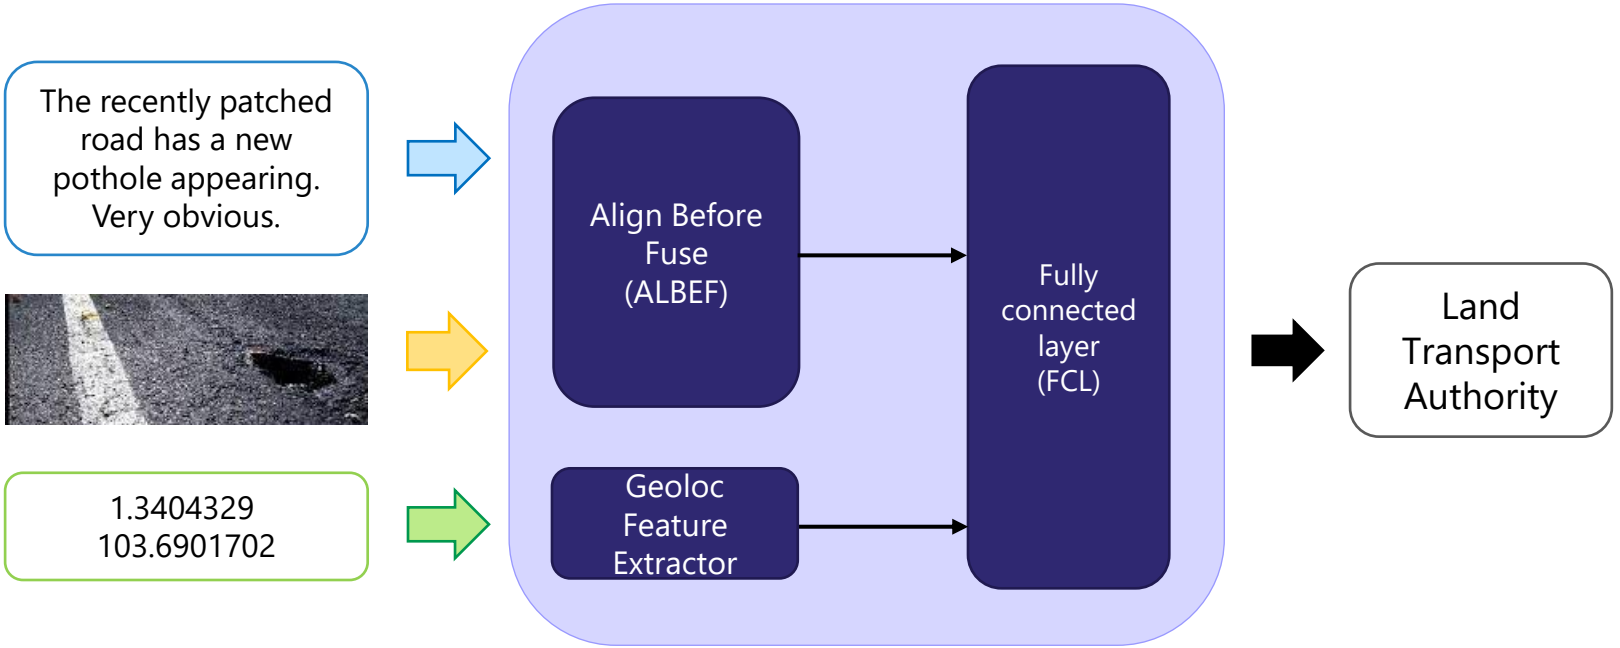

Geolocation Features

- Geolocation data provides valuable information in a small subset of cases where multiple agencies handle similar cases in the same area

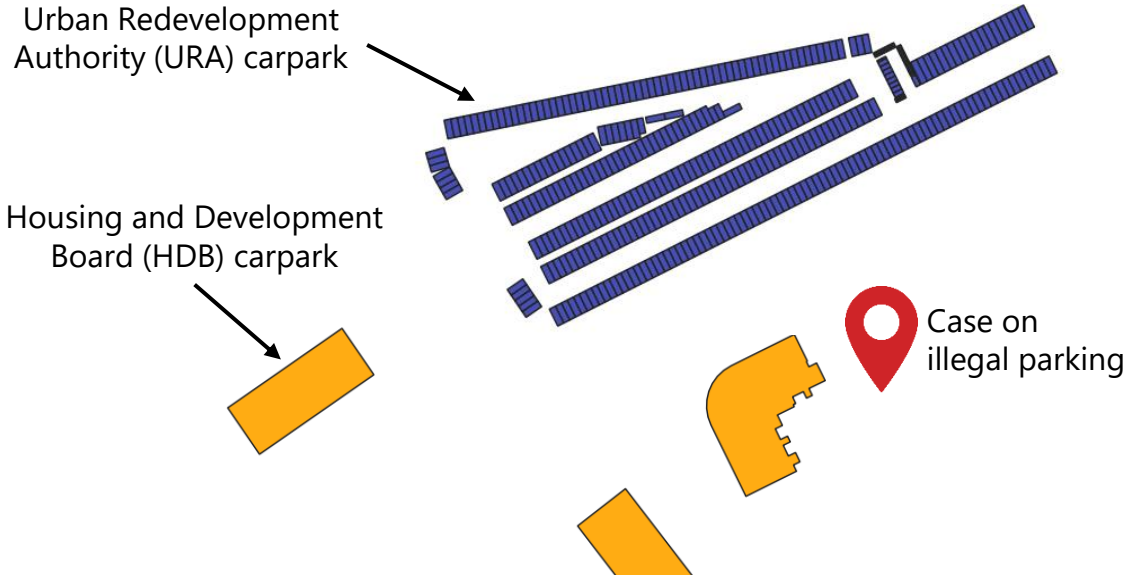

Both agencies own and maintain carparks in the same area.

Since the incident location is nearer to the orange areas, the case is likely under the purview of HDB.

Geolocation Feature Extraction

- Features are log distances to each map feature
- Output is a 14-dimension vector

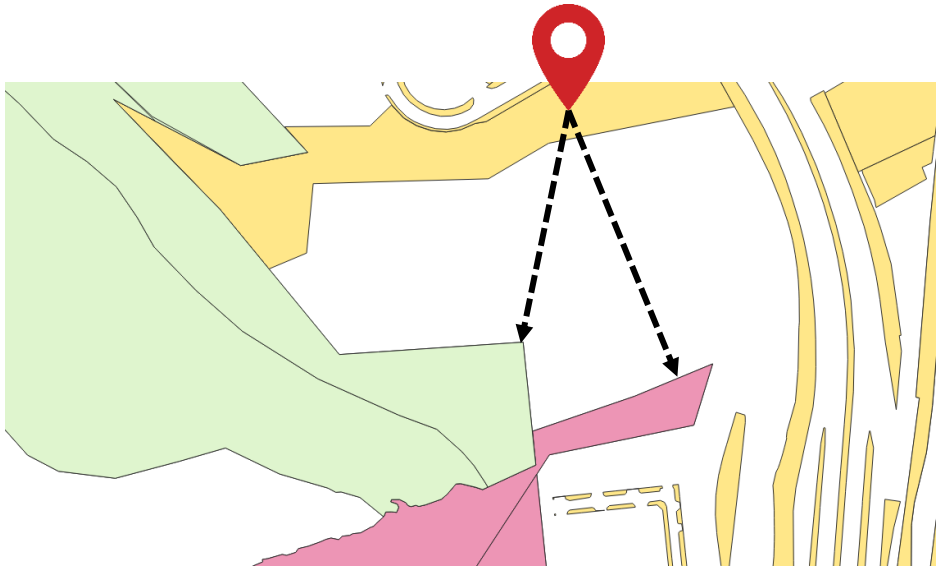

Land owner and land type	Log distance to nearest land parcel
HDB blocks	0.31
NParks Park connectors	1.00
SLA Nature reserves	0.65
…	…
JTC Market & Food Centres	0.00

Model Accuracy Comparison

Data used: 30k cases
Num classes: 13
Train-Test Split: 80%-20%

Randomly selected subset from the most recent data with images

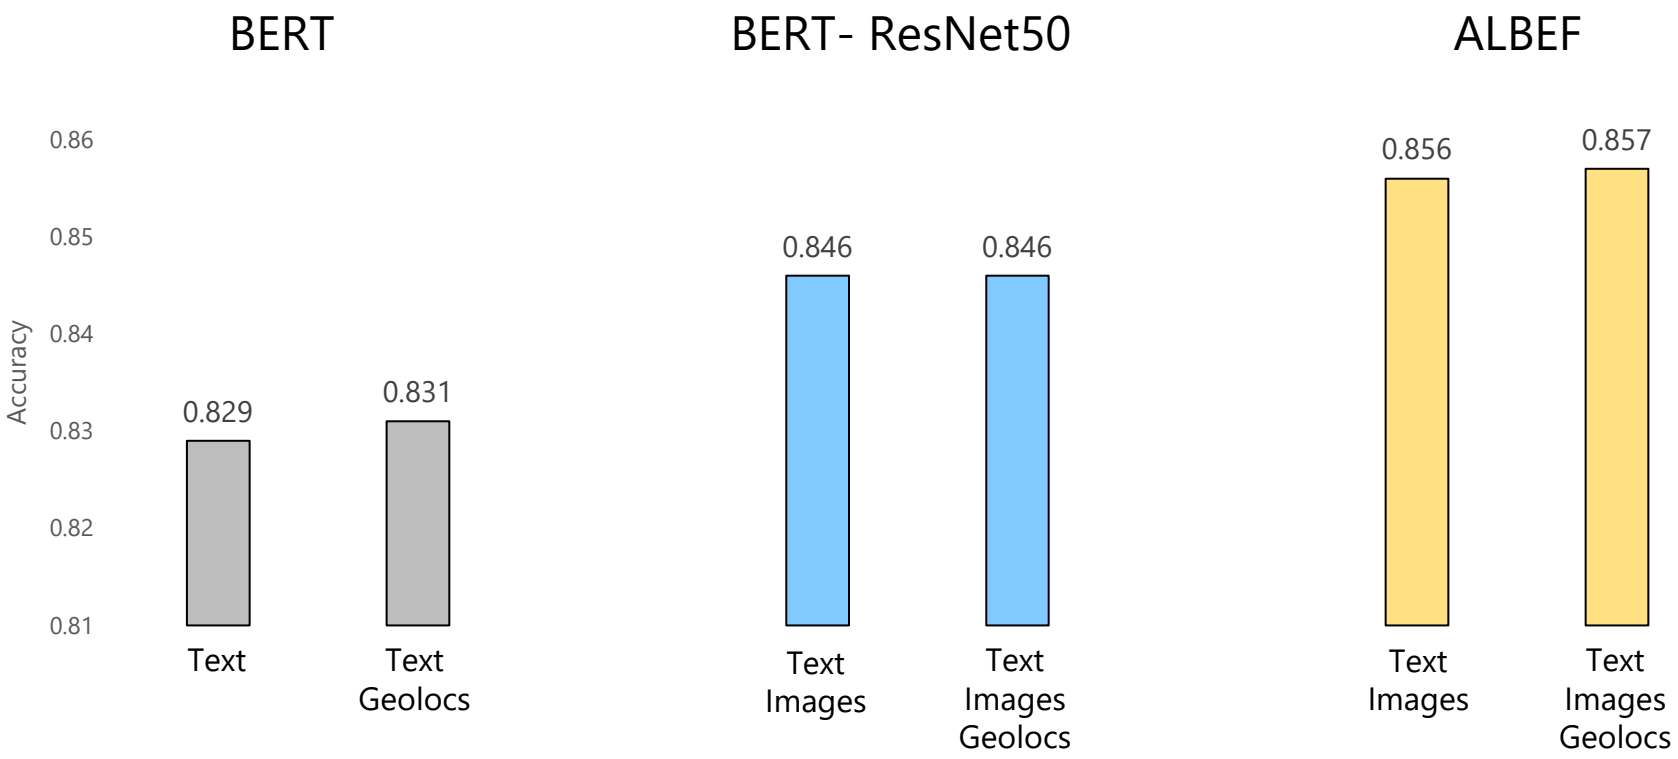

Hands-on Session

Objective (Recap)

- At the end of the tutorial, you will be able to use PyTorch and Hugging Face Transformers to build:
  - A text classifier using BERT
  - A text-image classifier with dual encoders using BERT and ResNet
  - A text-image classifier with joint encoders using ALBEF

- We exclude the geolocation portion from the tutorial because it is too specific to our use case

Dataset - WebVision

- The municipal issues dataset cannot be shared as it contains sensitive information
- WebVision dataset shares similar characteristics

- Text-image pairs are tagged with a class
- Images are crawled from the Flickr website and Google Images search
- Texts are captions, user tags or descriptions

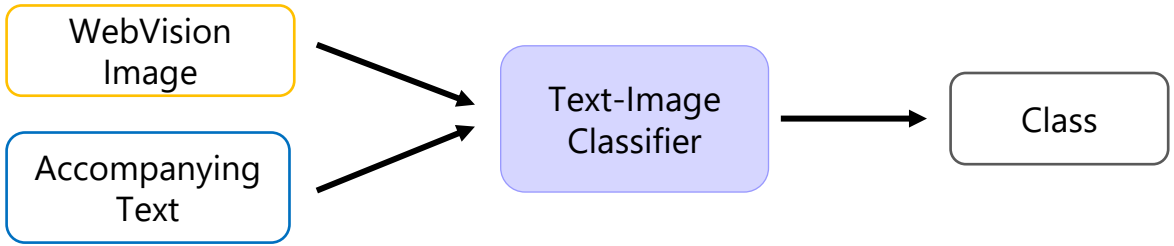

- We selected 1,000 cases from 10 classes for the hands-on session, with a 80%-20% train-test split

Sample Text-Image Pairs (Municipal Issue Dataset)

Text: The covered walkway outside Dakota Station Exit B has holes in it and leaks water on your head, please fix it.

Class: LTA

Text: The receipt is jammed up and I could not retrieve a copy of the statement. Can someone clear up the jammed paper?

Class: HDB

Text: Long time never clean with dead bird and oil stain.

Class: PUB

Text: Fallen tree next to rubbish collection center.

Class: NParks

Text: The ponding has been around for few months. It could be a potential mosquito breeding ground.

Class: NEA

Sample Text-Image Pairs (WebVision Dataset)

Text: Looks like a ferret!
Class: Polecat

Text: Sea Waves over Sand Beach
Class: Breakwater

Text: Female wearing a respirator posing indoors
Class: Gasmask

Text: Book shop Istanbul, Turkey
Class: Bookshop

Text: Fishing Photos, Blue Marin
Class: Gar

Text: Sweet Candy Store Flyer Template 2
Class: Confectionary

Task 1: Build a Text Classifier with BERT

BERT Architecture

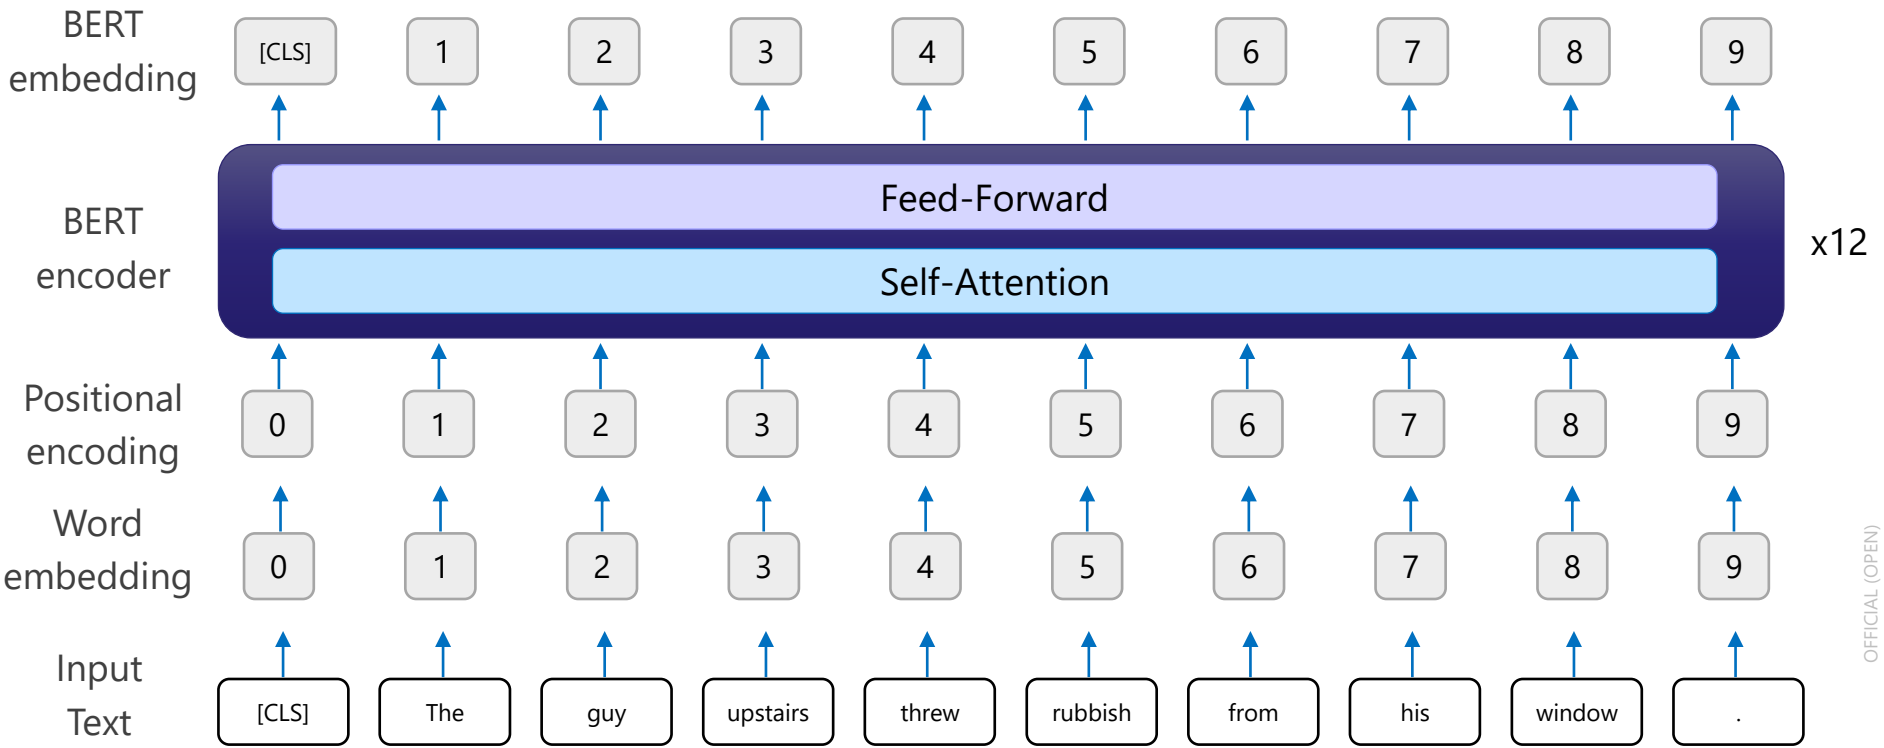

Text Encoder Architecture

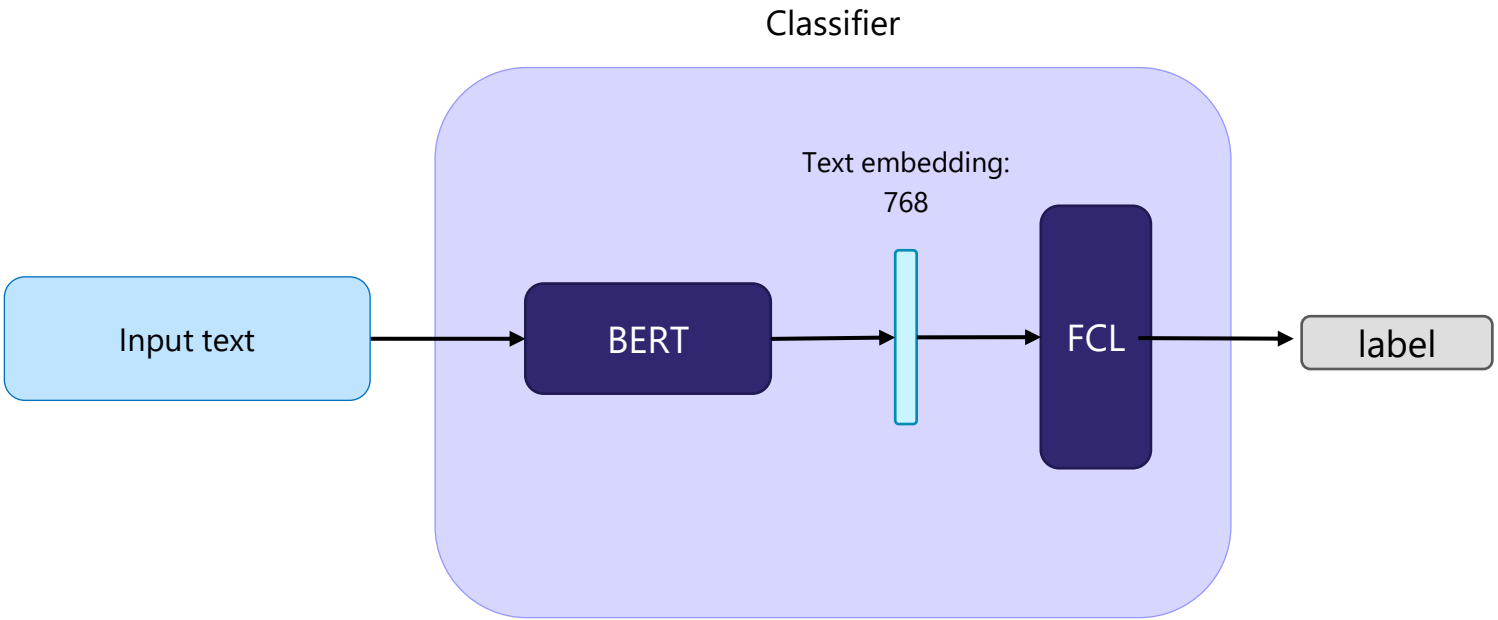

Task 2: Build a Dual-Encoder Text-Image Classifier with BERT and ResNet

Combining Text and Image

My little cutie is sleeping soundly on the floor.




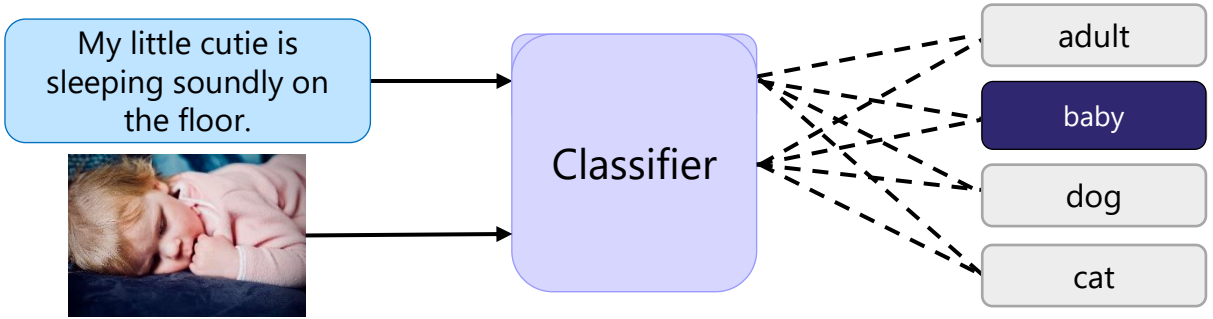

Dual-Encoder Text-Image Architecture

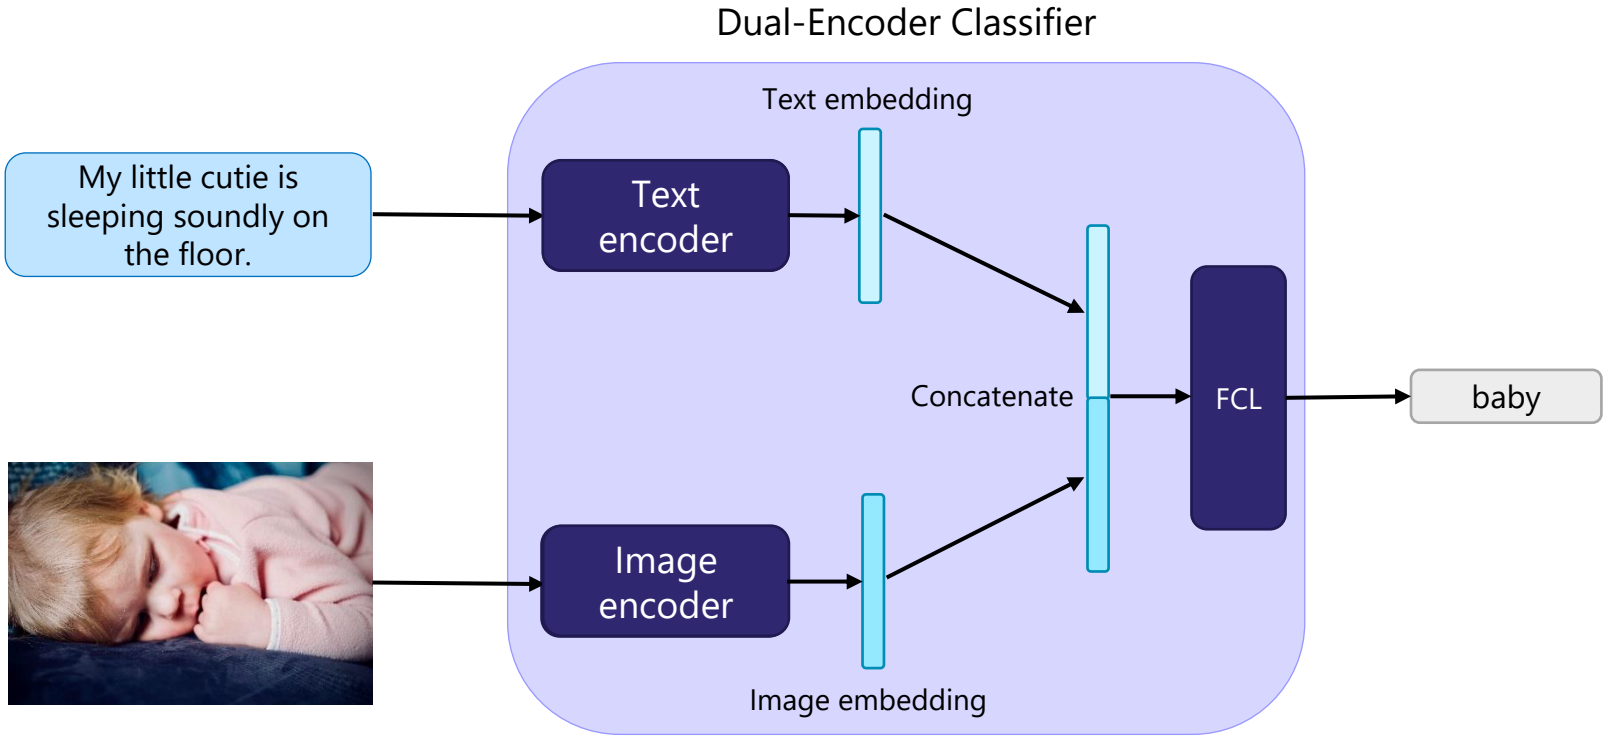

Image Encoder - ResNet

- Very deep network:
  - It can represent very complex functions
  - As the network increases in depth, classical CNNs do not perform well
- Solution: residual learning framework

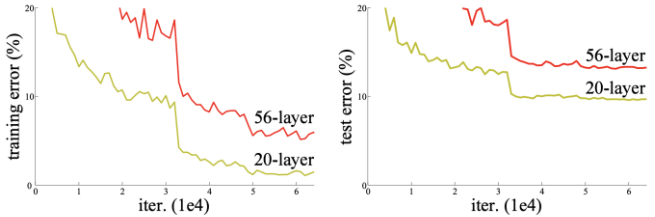

Figure 1. Training error (left) and test error (right) on CIFAR-10 with 20-layer and 56-layer "plain" networks. The deeper network has higher training error, and thus test error. Similar phenomena on ImageNet is presented in Fig. 4.

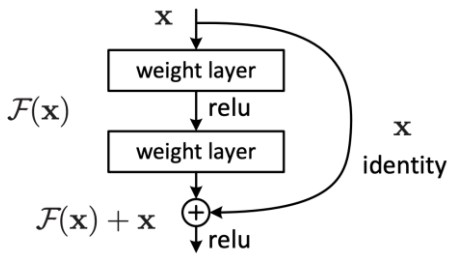

Figure 2. Residual learning: a building block.

"Deep residual learning for image recognition." He, et al. *CVPR*, 2016.

Image Encoder - ResNet50

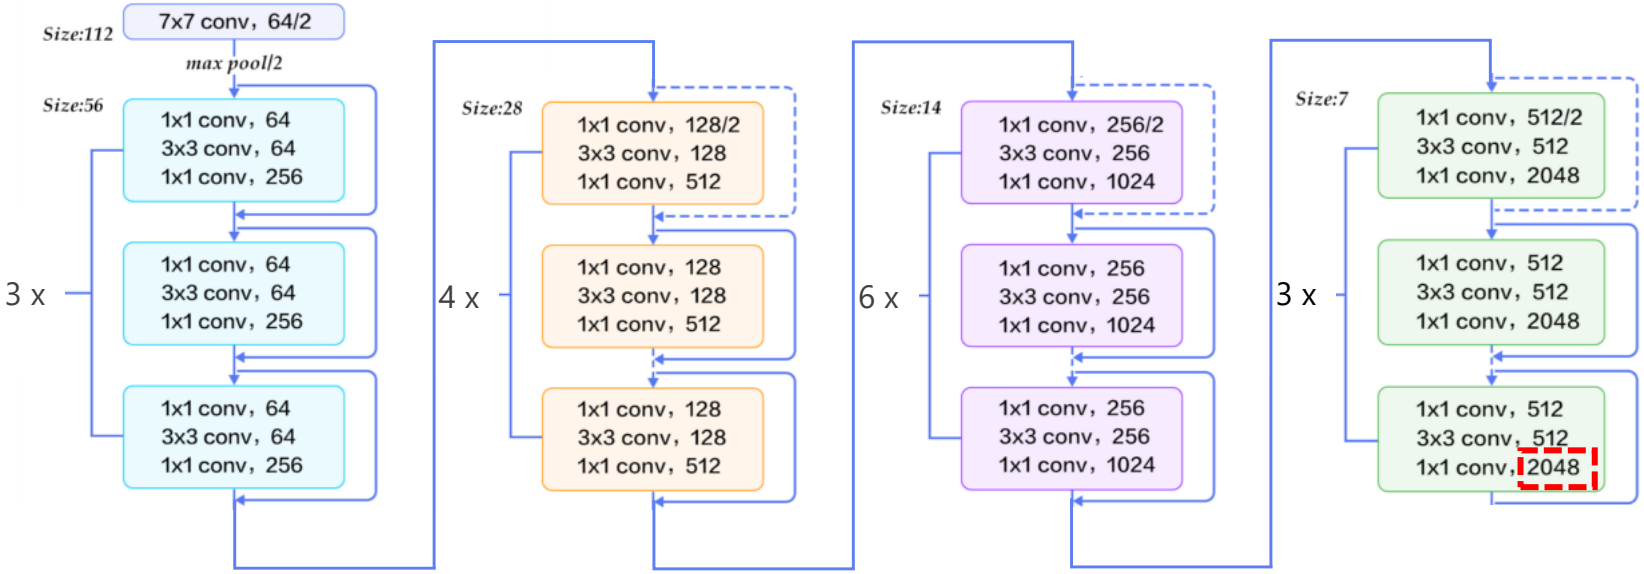

Dual-Encoder Text-Image Architecture

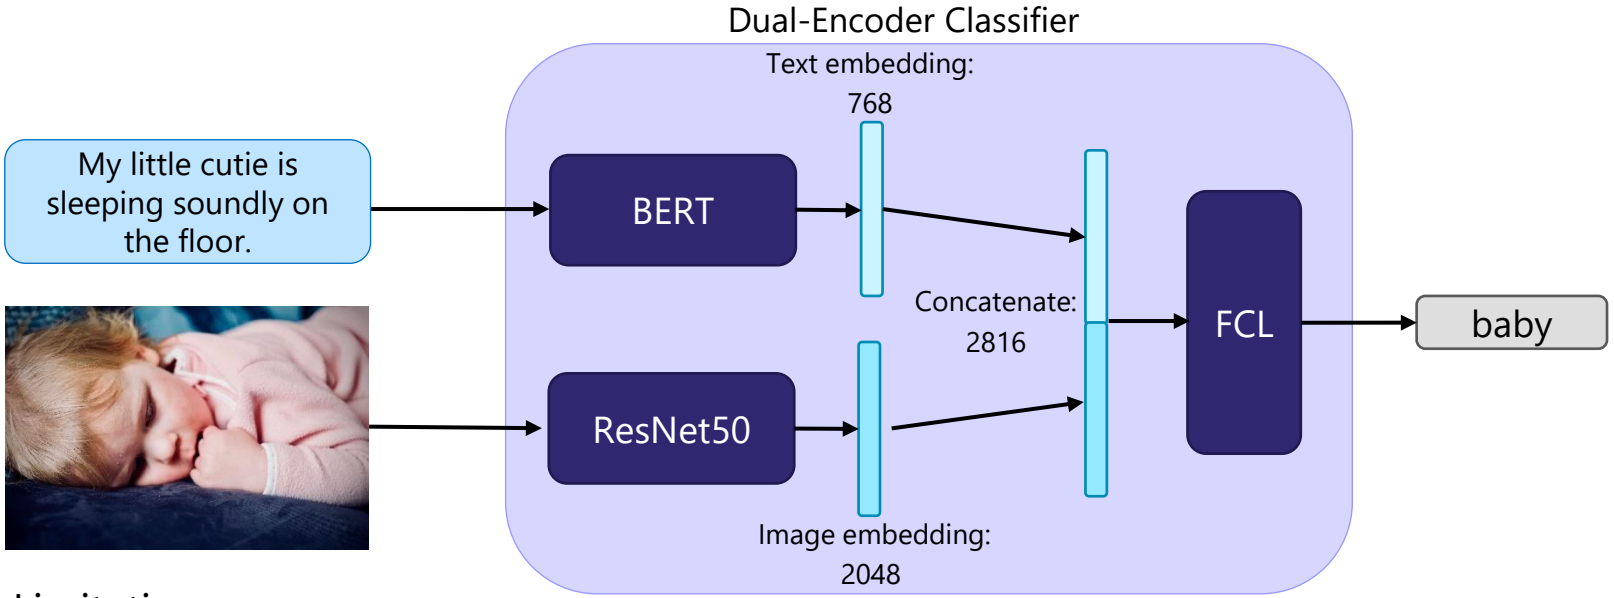

Limitations:

- Architecture only supports linear interactions between image and text
- Pretrained text and image encoders don't have knowledge about text-image pairs because they were pretrained separately on unimodal datasets

Task 3: Build a Joint-Encoder Text-Image Classifier with ALBEF

Dual-Encoder vs Joint-Encoder Architecture

Dual-Encoder Classifier

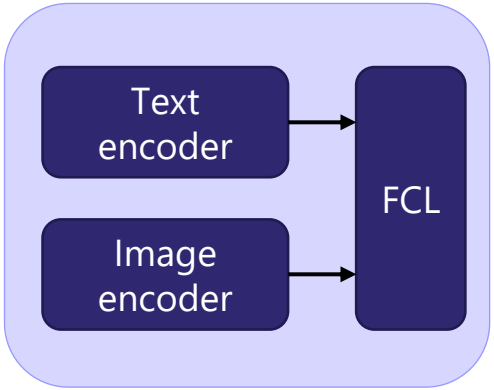

BERT + ResNet50

Joint-Encoder Classifier

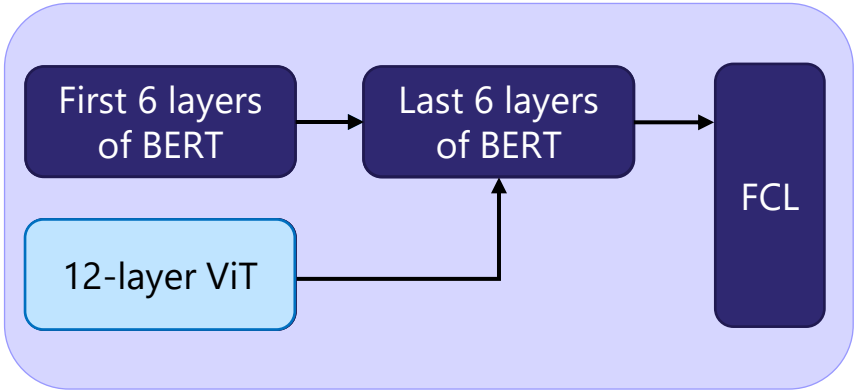

ALign the image and text representations BEfore Fusing (ALBEF)

"Align before fuse: Vision and language representation learning with momentum distillation." Li, et al. *NeurIPS*, 2021.

Image Encoder: Vision Transformer

- Vision Transformer (ViT) is the encoder network of transformer
- State-of-the-art for image classification

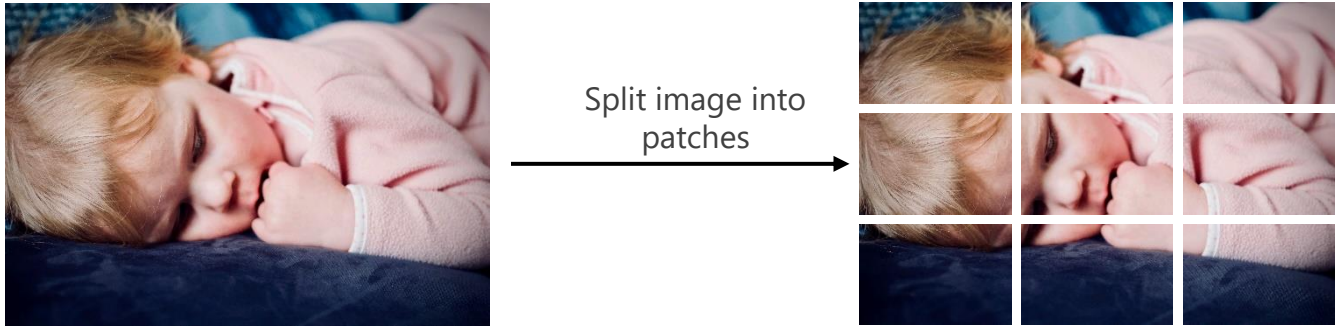

Image Encoder: Vision Transformer

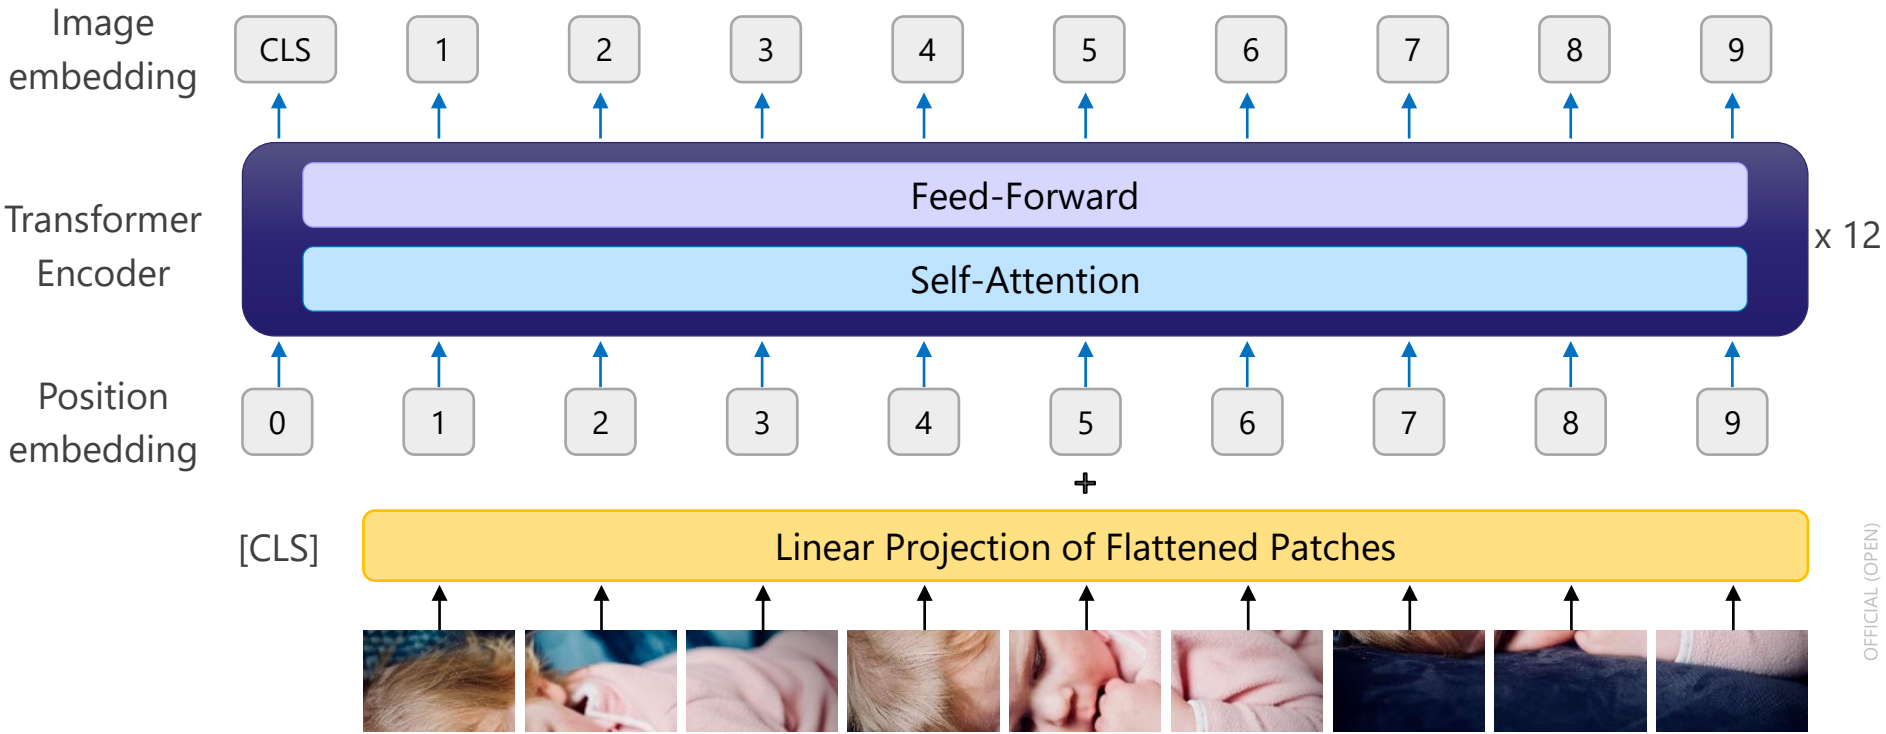

"An image is worth 16x16 words: Transformers for image recognition at scale." Dosovitskiy, et al. *arXiv:2010.11929*, 2020.

Joint-Encoder Text-Image: ALBEF

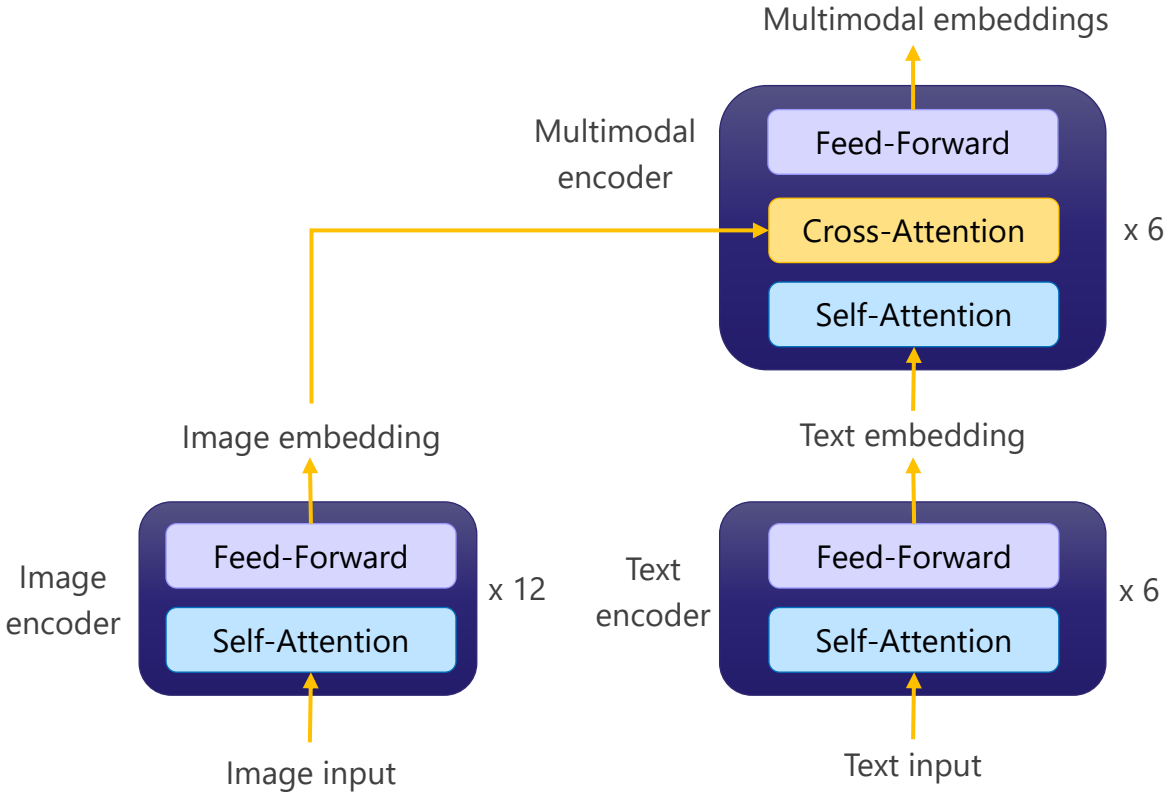

(Recap) limitations of dual-encoder:

- Architecture only supports linear interactions between image and text
- Pretrained text and image encoders don't have knowledge about text-image pairs because they were pretrained separately on unimodal datasets

ALBEF Pretraining - Data

Text-image pairs

Man sits in a rusted car buried in the sand on Waitarere beach

Little girl and her dog in northern Thailand. They both seemed interested in what we were doing

Interior design of modern white and brown living room furniture against white wall with a lamp hanging.

Emma in her hat looking super cute

ALBEF Pretraining - Objectives

- Image-Text Contrastive Loss: align the unimodal features

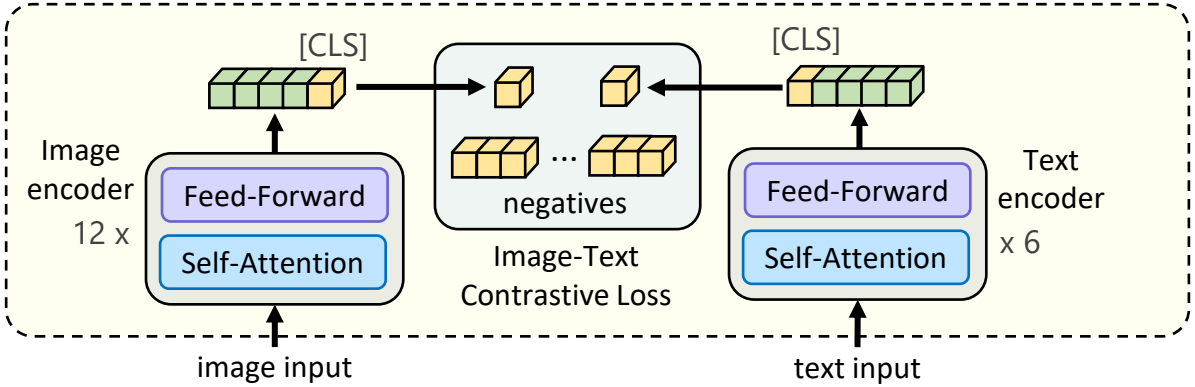

"Align before fuse: Vision and language representation learning with momentum distillation." Li, et al. NeurIPS, 2021.

ALBEF Pretraining - Objectives

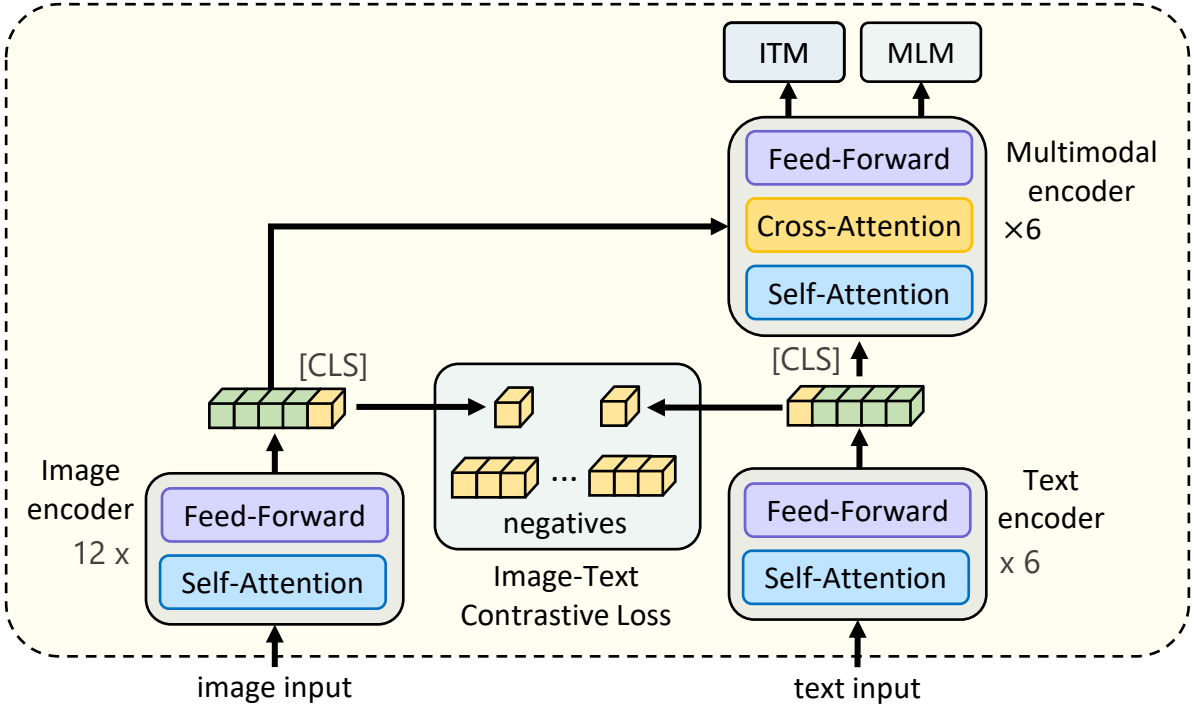

- Image-Text Contrastive Loss: align the unimodal features
- Image-Text Matching (ITM): binary classification of positives and negatives
- Mask Language Modeling (MLM): predict words based on image and contextual text

Joint-Encoder Text-Image Architecture

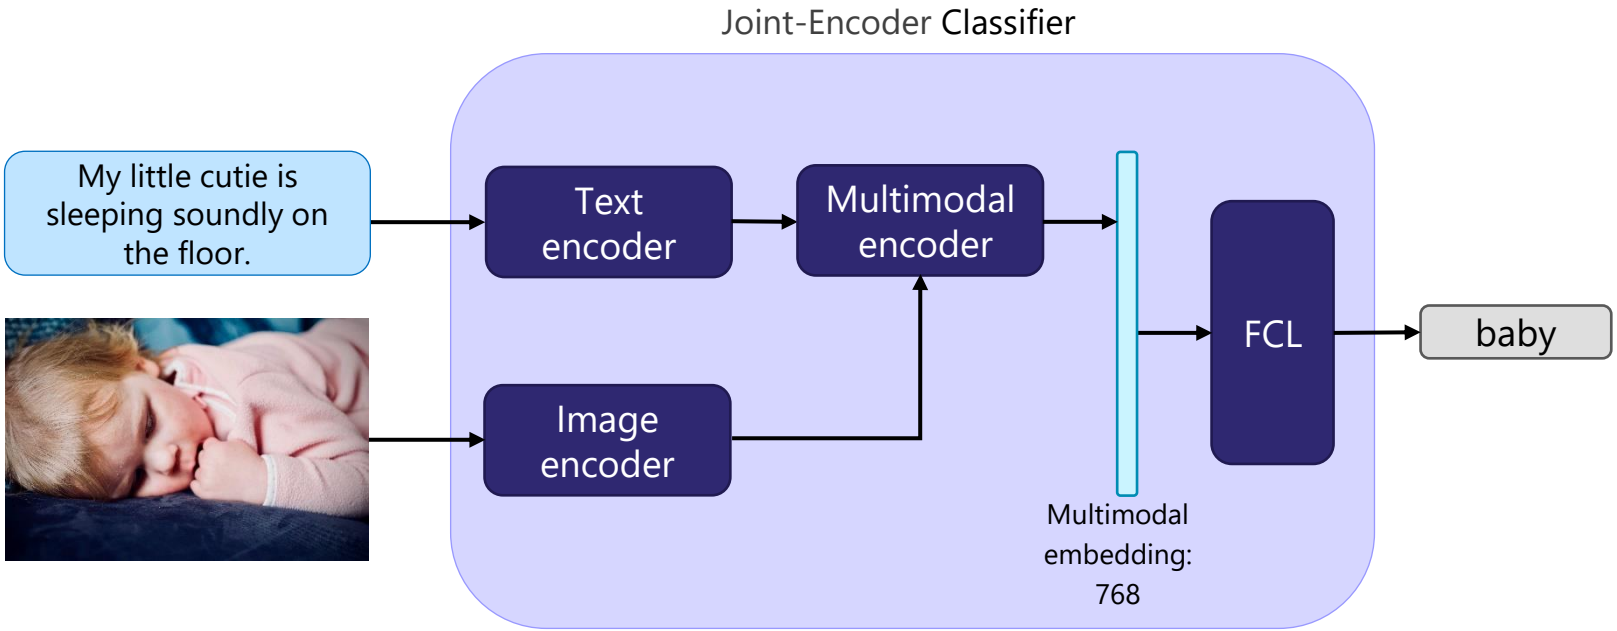

Discussion

Best Model Architecture?

- We have shown you **how** to train models using the different architectures on a toy dataset

- Proper evaluation on your own dataset is still required to pick the best model
  What is best for our dataset may not be best for yours

- Beyond model accuracy/per-class F1 scores, you should also consider:
  - Training time
  - Inference time
  - Resources required: GPU/RAM

Next Steps

- Hyperparameter tuning was not covered in this tutorial.
  You can do it on your own using these steps:

  - Split the dataset into train, validation and test sets
  - **Train set:** run *x (e.g. 50)* trials with different combinations of l*earning rate, learning rate scheduler, batch size, weight decay, training epochs*
  - **Validation set:** select the best model for each model architecture
  - **Test set:** evaluate the performance of the best hyperparameter tuned models for each model architecture, and select the best overall model

  - You can use tools like Optuna or Ray Tune
  - A good guide to hyperparameter tuning can be found here